

Anurag Thakur

+91 8626979168 | thakurmnnu011@gmail.com | LinkedIn - anuragthakur-660683241 | GitHub - NinjaKnight05
anuragthakur07.vercel.app

PROFESSIONAL SUMMARY

AI/ML Engineer specializing in Generative AI, RAG pipelines, and multi-agent systems using LangChain, CrewAI, and vector databases. Skilled in Python, prompt engineering, and workflow automation (n8n, Zapier), with a focus on production-ready, enterprise-grade AI solutions.

TECHNICAL SKILLS

Languages: Python, SQL, JavaScript, HTML, CSS, C

Generative AI: LLMs, RAG, LangChain, CrewAI, Langsmith, Prompt Engineering, NLP

Vector Databases: FAISS, Pinecone, ChromaDB, Qdrant, PostgreSQL, Firebase

ML & DL: Scikit-Learn, TensorFlow, PyTorch, OpenCV, Pandas, NumPy, PowerBI, Excel

Backend & APIs: Flask, Streamlit, RESTful APIs, FastAPI

Workflow Automation: n8n, Zapier

Cloud & DevOps: AWS (EC2, S3), Docker, Git, CI/CD Pipelines, DVC

EXPERIENCE

AI/ML Engineer Intern

Nov 2025 – May 2026

Coder Roots

- Designed and deployed 4+ production-grade ML pipelines for classification, regression, and clustering tasks using Scikit-Learn and XGBoost, applying rigorous evaluation frameworks and performance optimization.
- Built and deployed RESTful APIs for ML inference services with comprehensive documentation, enabling seamless integration with downstream systems and real-time predictions.
- Trained CNN and transfer-learning models using TensorFlow and PyTorch, achieving robust validation accuracy while managing multiple project priorities independently.

PROJECTS

Kodexa

Repo, Live

- Architected an autonomous AI agent that indexes GitHub repositories and answers developer queries via RAG, demonstrating expertise in agentic AI design and tool-using agents.
- Integrated Pinecone vector database, LangChain, and Groq Llama 3 for semantic code search, achieving sub-2s query latency with optimized retrieval pipelines.
- Designed REST endpoints for repository indexing and concurrent multi-repo session management, supporting real-time developer interactions.

PaperPilot AI

- A full-stack RAG study assistant (FastAPI + React) that pulls answers from an uploaded PDF, the web, or general knowledge, and labels which source each answer came from.
- PDF upload with automatic text chunking and retrieval, plus a React chat UI with per-mode history and delete/switch controls.

Energy Consumption Forecasting

Repo, Live

- Developed an XGBoost-based forecasting application on 121,000+ hourly records from PJM grid data, demonstrating strong data analysis and feature engineering skills.
- Engineered time-based features (lag windows, hour, day, seasonality) improving R^2 from 0.45 to 0.87 through iterative optimization.
- Deployed an interactive Streamlit dashboard with date-range selection, peak-consumption metrics, and trend-analysis charts for real-time data exploration.

Driver Drowsiness Alert System

Repo

- Built a real-time drowsiness detection system using OpenCV and dlib for live video feed processing with sub-2s alert latency.

- Applied Eye Aspect Ratio (EAR) thresholding to distinguish prolonged eye closure, implementing guardrails to minimize false positives.

EDUCATION

CCPA: 6.57

DAV University: Btech - Computer Science Engineering.

- Participated in an inter-college technical hackathon (CGC Jhanjeri, Nov 2023), building a working prototype within a 24-hour constraint, demonstrating curiosity, initiative, and a commitment to learning emerging technologies.